

Assignment 1

In this assignment you will be learning about and implementing [Markov Chains](#), if you have never heard of them before, do not worry, this notebook will explain everything needed for the assignment.

Markov chains

The best way to introduce what a Markov chain is with a simple example. Suppose you have a simple counter that you can set to any number between 0 and 9 (including 0 and 9). At every step you roll a standard 6 -sided die and depending on whether the rolled number is even or odd you either increment or decrement the counter. If you try to increment at 9, lets say it wraps around to 0 and similarly for trying to decrement at 0. Now you can play this game for multiple steps and record the counter value after every step, this is just a simple markov chain.

Here is an example evolution of the markov chain, Start with counter at 1.

- Step 1 : Roll a 4, its even so increment the counter to 2.
- Step 2 : Roll a 2, its even so increment the counter to 3.
- Step 3 : Roll a 5, its odd so decrement the counter to 2.

etc.

A markov chain consists of a set of n states (in our example it was the 10 states of the counter) and a probabilistic rule to jump to another state every step (the dice rolls). Although in our simple example we had a similar rule for each state, you can have a different rule for each state. For example you may say that if you the counter is odd, dont roll a die, always increment, and follow the die for even states. The rule is always in the form of n probabilities that sum to one. Each probability indicates the chance that it jumps to that state.

In our example the rule for state 0 can be represented by the below numpy vector

When you are in state 0, you have 50% chance of landing on state 1 (you rolled even) or a 50% chance of landing in state 9 (you rolled odd and tried to decrement at 0)

```
import numpy as np
np.array([0. , 0.5, 0. , 0. , 0. , 0. , 0. , 0. , 0. , 0.5])
```

The index in the array represents the state and the value represents the probability with which you will jump to that state. Similarly this is the rule for state 4

```
np.array([0. , 0. , 0. , 0.5, 0. , 0.5, 0. , 0. , 0. , 0. ])
```

Transition Matrix

The way to represent markov chains is with a transition matrix \mathbf{T} of size (n,n) , where the column i will be the rule vector for the state i (Here we are assuming that the states are numbered 0 to $n-1$).

Thus you can interpret the element T_{ij} of the matrix as the probability you will jump to state i from the state j .

Question 1

Write a function to return the transition matrix for the following markov chain, given the number of states as an argument.

The rule for the markov chain is as follows,

If current state is i

1. With probability 0.8, it goes to state $i+1$
2. With probability 0.2, it will fall back to state $i = 0$.
3. At state $n-1$, since it cannot go higher, with probability 1 it stays in the same state of $n-1$.

Intuitively, this markov chain behaves in the following way, at every step there is an 80% chance that it climbs up by 1 and 20% chance that it falls down completely, however if it does reach the final state $n-1$, it stays there forever

You are **not** allowed to use any for loops for this question, instead use numpy indexing to fill in your matrix. You can assume that n is atleast 3.

Hints:

1. First write your code in an empty cell for with a hardcoded n value, in this way you can quickly iterate and get the right answer first. Then copy it into the function.
2. You can use [Array indexing](#)

```
#YOUR CODE HERE
def transition_matrix(n):
    pass

#TESTING
transition_matrix(10)
```

To test your function , check that the above returns

```
array([[0.2, 0.2, 0.2, 0.2, 0.2, 0.2, 0.2, 0.2, 0.2, 0. ],
       [0.8, 0. , 0. , 0. , 0. , 0. , 0. , 0. , 0. , 0. ],
       [0. , 0.8, 0. , 0. , 0. , 0. , 0. , 0. , 0. , 0. ],
       [0. , 0. , 0.8, 0. , 0. , 0. , 0. , 0. , 0. , 0. ],
       [0. , 0. , 0. , 0.8, 0. , 0. , 0. , 0. , 0. , 0. ],
       [0. , 0. , 0. , 0. , 0.8, 0. , 0. , 0. , 0. , 0. ],
       [0. , 0. , 0. , 0. , 0. , 0.8, 0. , 0. , 0. , 0. ],
       [0. , 0. , 0. , 0. , 0. , 0. , 0.8, 0. , 0. , 0. ],
       [0. , 0. , 0. , 0. , 0. , 0. , 0. , 0.8, 0. , 0. ],
       [0. , 0. , 0. , 0. , 0. , 0. , 0. , 0. , 0.8, 0. ]],
      dtype=float64)
```

```
[0. , 0. , 0. , 0. , 0. , 0. , 0. , 0. , 0.8, 0. , 0. ],
[0. , 0. , 0. , 0. , 0. , 0. , 0. , 0. , 0. , 0.8, 1. ]])
```

Question 2

Now we want to use the transition matrix to know what are the probabilities of being in each state after running the Markov chain for k steps. For example, for the above Markov chain, I want to know after 30 steps, what is the probability that I have reached the state $n-1$.

Luckily this is fairly easy to calculate, using matrix multiplication. An amazingly useful property of the transition matrix is as follows, if you have vector p_j as the vector of probabilities after step j , then p_{j+1} is given by the following matrix vector product.

$$p_{j+1} = T p_j$$

We will not prove this result here, but if you spend some time thinking about what matrix multiplication is doing here, especially interpret it as taking a linear combination of columns then it should make sense.

Thus if someone gave you a probability distribution for the starting state as p_0 , then the probability distribution after k steps is simply given by multiplying by the transition matrix k times

$$p_k = T^k p_0$$

Implement a function which given an initial distribution p_0 , number of steps k and a transition matrix tm , it will return the distribution after k steps

```
#YOUR CODE HERE
def propagate(p0,k,tm):
    pass
#TESTING
tm = transition_matrix(10)
p0 = np.zeros(10)
p0[0] = 1
pk = propagate(p0,30,tm)
pk
```

The above code should return :

```
array([0.0816156 , 0.06787354, 0.05644353, 0.04694437, 0.03904617,
       0.03247669, 0.02701094, 0.02246264, 0.01867741, 0.6074491 ])
```

Question 3

Now use the above function to plot the probabilities at the end of each step for the first 10 steps in the same plot.

1. Create a transition matrix for $n = 10$

2. Create a vector which represents the initial state that you are in state 0 with probability 1.
3. Plot the probability distribution after each step for 10 steps in the same figure, You will have one curve for each time step, so a total of 10 curves in one plot
4. Make sure you label the X and Y axis, and give the plot a title, save the plot in a file called "qsn3.png". You can do so with the following code `plt.savefig("qsn3.png")`

#YOUR CODE HERE

Question 4

Part a

Use the functions you implemented to calculate the following :

How many steps does it take for the probability of being in the final state to be atleast 0.5.

After each step the probability of being in the final state slowly increases. At some point this probability will cross 50%. What you need to calculate is how many steps does it take to cross 50%.

Again use $n=10$ transition matrix and start from state 0.

Hints:

1. Use the while loop

#YOUR CODE HERE

Part b

Convert the above Code into a function that computes the number of steps to 50% probability in the final state given n the size of the markov chain

```
#YOUR CODE HERE
def num_steps(n):
    pass
```

Part c

Compute and plot the number of steps required for n ranging from 10 to 40.

Lookup the plotting function `plt.semilogy` and understand what it does, plot another graph of the number of steps using it.

What does the semilogy plot tell you?

Save the plots as "qsn4c.png" and "qsn4c_semilogy.png"

#YOUR CODE HERE

Question 5

Now we want to also compute some sample evolutions of the markov chains. What does this mean? We want a sequence of states drawn according to the rules of the markov chain. Thus we want to use randomness to choose what the next state is. Thus each call to the function will return a different answer, a new sample.

For example consider the markov chain we have been using, it climbs up with probability 80% or drops to 0 otherwise, once it reaches the last state, it is stuck there. Here is one such sample of 20 steps.

```
[0, 1, 2, 3, 0, 1, 2, 3, 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 9, 9]
```

Write a function that returns one such sample for each call. It should take the transition matrix `tm`, number of steps `k` and a starting state `s0` and return an array of size `k+1` of states (including the starting state).

```
import numpy.random as rnd
#YOUR CODE HERE
def sample(tm, k, s0):
    pass
```

Use the below code to test out some of the samples, it samples 20 trajectories of the markov chain and plots them. The plot should show that almost all or all the trajectories would have reached the final state by the end. Save this graph as "qsn5.png".

```
#TESTING
import matplotlib.pyplot as plt
tm = transition_matrix(10)
k = 20
s0 = 0

for i in range(20):
    states = sample(tm, 100, s0)
    plt.plot(states)

plt.savefig("qsn5.png")
```

Question 6

Since the states are roughly being incremented by 1 or falling to 0, we want to know what rate of progress looks like on average.

1. Create a transition matrix for `n = 25`
2. Sample 1000 trajectories of 100 steps using the sample function you wrote above.
3. Compute the average state (over the 1000 samples) at each time step.
4. Plot the average state vs time steps and save it as "qsn6.png"

Hints:

1. You should see a plot where the average climbs quickly and then steadily increases slowly till the end

#YOUR CODE HERE

Question 7

We want to now verify if our sampling implementation is consistent with our implementation of computing the probability distribution.

1. Create a transition matrix for $n = 25$
2. Sample 1000 trajectories of 100 steps using the sample function you wrote above, record the last state for each trajectory in a separate list.
3. Plot a histogram of final states
4. Use the function `propagate` that you wrote to compute the theoretical distribution after 100 steps
5. Plot the expected distribution in the same plot and verify that it matches with the histogram.
6. Save the plot as "qsn7.png"

Hints:

1. For the the histogram you can set the bins manually to be one bin for each state. You can do this by setting bin boundaries to be ... 2.5, 3.5, ... , in this way the state 3 will have a bin between 2.5 and 3.5
2. The `propagate` functions returns a probability distribution, you need to multiply it with the number of samples for it to be comparable to the histogram

#YOUR CODE HERE

Submission Instructions (we talk in class)